

Program Structure and Algorithms

Sid Nath

Lecture 5

Agenda


- Administrative
 - HW2 questions?
 - Quiz 3 today on D/Q
- Lecture
- Quiz

Graph Applications

- Numerous
- Networks: social, transportation, circuits, ...
- Internet
- Maps
- OS
- Backprop in neural networks
-

Graph Definition

- A *graph* $G = (V, E)$ consists of two things:
 - A collection V of *vertices*, or objects to be connected.
 - A collection E of *edges*, each of which connects a pair of vertices.
 - May be undirected or directed
- Examples
 - The internet $V = \{\text{websites}\}$, $E = \{\text{links}\}$ or $V = \{\text{computers}\}$, $E = \{\text{physical connections}\}$
 - Highway system, $V = \{\text{intersections}\}$, $E = \{\text{roads}\}$



$V = \{A, B, C, D, E\}$
 $E = \{AB, AC, AD, BD, CE, DE\}$

Graph Representations

How do you store a graph in a computer?

- **Adjacency matrix:** Store list of vertices and an array $A[i,j] = 1$ if edge between v_i and v_j .
 - Small space for dense graphs.
 - Slow for most operations.
- **Edge list:** List of all vertices, list of all edges
 - Hard to determine edges out of single vertex.
- **Adjacency list:** For each vertex store list of neighbors.
 - Needed for DFS to be efficient
 - We will usually assume this representation

How to Represent Graphs

- Adjacency matrix

```
0
 / \
1---2
 \ /
 3
```

```
# Create a graph with 4 vertices and 5 edges
graph = [[0, 1, 1, 0],
         [1, 0, 1, 1],
         [1, 1, 0, 1],
         [0, 1, 1, 0]]
```

- Adjacency list

```
# Create a graph with 4 vertices and 5 edges
graph = {0: [1, 2],
        1: [0, 2, 3],
        2: [0, 1, 3],
        3: [1, 2]}
```


- Edge list

Graph Operations

- Search
 - Vertex reachability
 - Enter/exit times
 - Component the vertex belongs to

Basic Algorithm

Keep track of all areas discovered
While there is an unexplored path,
follow path



Systematize

Need to keep track of:

- Which vertices discovered
- Which edges have yet to be explored

Explore Algorithm will:

- Use a field `v.visited` to let us know which vertices we have seen.
- Store edges to be explored implicitly in the program stack.

Explore

```
explore(v)
```

```
  v.visited  $\leftarrow$  true
```

```
  For each edge (v,w)
```


```
    If not w.visited
```

```
      explore(w)
```

```
      w.prev  $\leftarrow$  v
```

Example


Note: edges used leave behind “DFS tree”.



```
explore(A)
  explore(B)
    explore(A)
      explore(C)
        explore(A)
          explore(B)
            explore(C)
              explore(D)
                explore(A)
                  explore(H)
                    explore(D)
                      explore(F)
                        explore(E)
                          explore(A)
                            explore(F)
                              explore(G)
                                explore(F)
                                  explore(H)
                                    explore(E)
```

Question: explore

Which vertices does `explore(v)` mark as visited?



Depth First Search

`explore` only finds the part of the graph reachable from a single vertex. If you want to discover the entire graph, you may need to run it multiple times.


```
DepthFirstSearch(G)
```

```
    Mark all  $v \in G$  as unvisited
```

```
    For  $v \in G$ 
```

```
        If not  $v.visited$ , explore(v)
```

Example



`explore(A)`
`explore(D)`

DFS(G) eventually discovers all vertices in G.

Runtime of DFS

```
explore(v)
  v.visited ← true
  For each edge (v, w)
    If not w.visited
      explore(w)
```

Run once per vertex $O(|V|)$ total

Run once per neighboring vertex $O(|E|)$ total

```
DFS(G)
  Mark all  $v \in G$  as unvisited
  For  $v \in G$ 
    If not v.visited, explore(v)
```

$O(|V|)$

Final runtime: $O(|V| + |E|)$

Note on Graph Algorithm Runtimes

Graph algorithm runtimes depend on both $|V|$ and $|E|$. (Note $O(|V|+|E|)$ is linear time)

What algorithm is better may depend on relative sizes of these parameters.

Sparse Graphs:

$|E|$ small ($\approx V$)

Examples:

- Internet
- Road maps

Dense Graphs:

$|E|$ large ($\approx V^2$)

Examples:


- Flight maps
- Wireless networks

Connected Components

- Want to understand which vertices are reachable from which others in graph G .
- `explore(v)` finds which vertices are reachable from a given vertex.

Theorem: The vertices of a graph G can be partitioned into *connected components* so that v is reachable from w if and only if they are in the same connected component.

Example



Problem: Computing Connected Components

Given a graph G , compute its connected components.

Run `explore(v)` to find the component of v . Repeat on unclassified vertices.

DFS lets us do this!

ConnectedComponents (G)

CCNum \leftarrow 0

For $v \in G$

$v.visited \leftarrow false$

For $v \in G$

 If not $v.visited$

 CCNum++

 explore(v)

explore(v)

$v.visited \leftarrow true$

$v.CC \leftarrow CCNum$

 For each edge (v,w)



 If not $w.visited$

 explore(w)

Runtime $O(|V|+|E|)$.

Example

CCNum: 2



Discussion about DFS

What does DFS actually *do*?

- No output.
- Marks all vertices as visited.

DFS also is a useful way to explore the graph.

By *augmenting* the algorithm a bit (like we did with the connected components algorithm), we can learn useful things.

Pre- and Post- Orders

- Keep track of what algorithm does & in what order.
- Have a “clock” and note time whenever:
 - Algorithm visits a new vertex for the first time.
 - Algorithm finishes processing a vertex.
- Record values as `v.pre` and `v.post`.

Computing Pre- & Post- Orders

DFS (G)

clock $\leftarrow 1$

For $v \in G$

$v.\text{visited} \leftarrow \text{false}$

For $v \in G$

 If not $v.\text{visited}$

 explore(v)

explore(v)

$v.\text{visited} \leftarrow \text{true}$

$v.\text{pre} \leftarrow \text{clock}$

 clock++

 For each edge (v, w)

 If not $w.\text{visited}$


 explore(w)

$v.\text{post} \leftarrow \text{clock}$

 clock++

Runtime $O(|V| + |E|)$.

Example



What do these orders tell us?

Prop: For vertices v, w consider intervals

$[v.pre, v.post]$ and $[w.pre, w.post]$.

These intervals:


1. Contain each other if v is an ancestor/descendant of w in the DFS tree.
2. Are disjoint if v and w are cousins in the DFS tree.
3. Never interleave
 $(v.pre < w.pre < v.post < w.post)$

Directed Graphs

Often an edge makes sense both ways, but sometimes streets are one directional.

Definition: A directed graph is a graph where each edge has a direction. Goes *from* v *to* w .


Draw edges with arrows to denote direction.





DFS on Directed Graphs

- Same code
- Only follow *directed* edges from v to w .
- Runtime still $O(|V|+|E|)$
- `explore(v)` discovers all vertices reachable from v following only directed edges.

Example





DFS on Directed Graphs



- A is *root* of tree; all other nodes are A's *descendants*
- E has *descendants* F, G, H (E is an *ancestor* of G)
- C is the *parent* of D
- H is a *child* of E

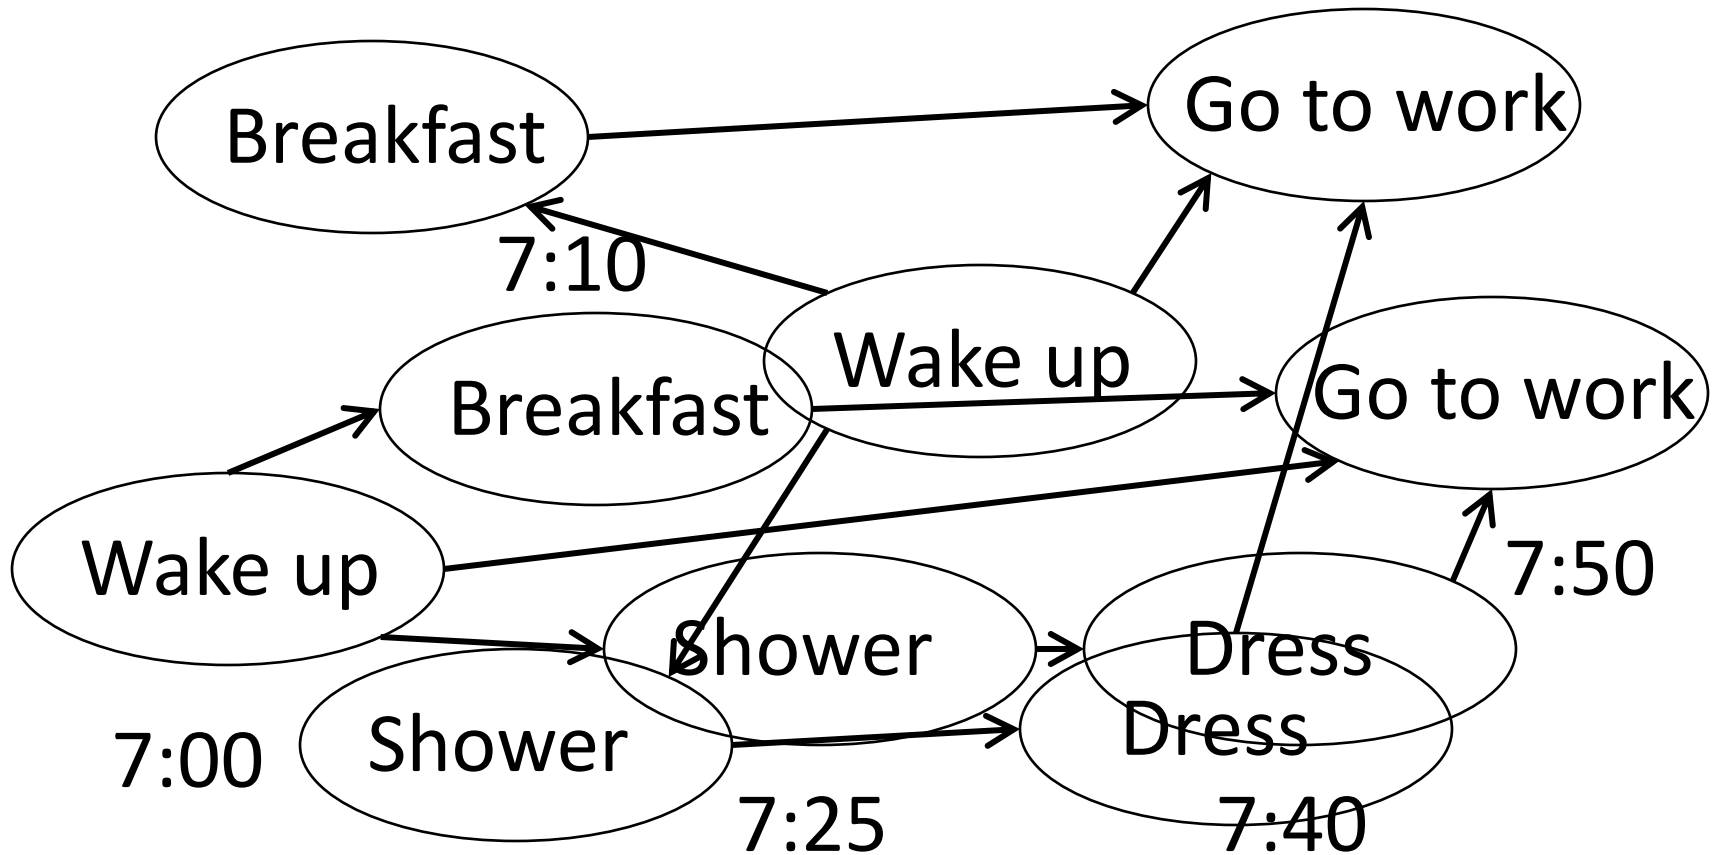
TERMINOLOGY

Ancestry and Pre/Post Numbers



- **Node u is an ancestor of node v iff $\text{pre}(u) < \text{pre}(v) < \text{post}(u)$**
- *Because u is an ancestor of v iff u is discovered first, and then v is encountered during the exploration of u*
- [Def. Node v is a *descendant* of u iff node u is an *ancestor* of v]

Dependency Graphs



Dependency Graphs


A directed graph can be thought of as a graph of dependencies. Where an edge $v \rightarrow w$ means that v should come before w .

Definition: A topological ordering of a directed graph is an ordering of the vertices so that for each edge (v,w) , v comes before w in the ordering.

Edge classification in Directed Graphs

- Forward edges are between a node and its non-child descendent
- Back edges lead to an ancestor node
- Cross edges lead to neither ancestor or descendent; they lead to a node that has already been completely explored (i.e., post visited)
- Tree edges are between parent and children.


DFS tree



pre/post ordering for (u, v)				Edge type
[[]]	Tree/forward
u	v	v	u	
[[]]	Back
v	u	u	v	
[]	[]	Cross
v	v	u	u	

Cycles

Definition: A cycle in a directed graph is a sequence of vertices $v_1, v_2, v_3, \dots, v_n$ so that there are edges $(v_1, v_2), (v_2, v_3), \dots, (v_{n-1}, v_n), (v_n, v_1)$



Cycles in Directed Graphs

- ***A directed graph has a cycle iff DFS reveals a back edge***
- (\Leftarrow) If (u,v) is a back edge, then it along with the $v \rightarrow u$ path in the search tree will form a cycle.
- (\Rightarrow) If the graph has a cycle $v_0 \rightarrow v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_k \rightarrow v_0$ then consider the node with smallest *pre* number, call it v_i . All other v_j on the cycle are reachable from v_i and will therefore be descendants of v_i in the search tree. Thus, the edge $v_{i-1} \rightarrow v_i$ is a back edge.
- So, we can determine whether G is acyclic in linear time


DAGs

Definition: A Directed Acyclic Graph (DAG) is a directed graph which contains no cycles.

Facts:

Let G be a (finite) DAG. Then G has a topological ordering.

Every finite DAG contains at least one sink.



Sources and Sinks in a DAG

- Node with the smallest post number has no outgoing edges.
 - Outdegree = 0
 - Node is called a sink
- Node with the largest post number has no incoming edges
 - Indegree = 0
 - Node is called a source

Topological Ordering and DAGs

- Prelude to shortest paths
- Generic scheduling problem
- Input:
 - Set of tasks $\{T_1, T_2, \dots, T_n\}$
 - Example: getting dressed in the morning: put on shirt, socks, shoes, ..
 - Set of dependencies $\{T_1 \rightarrow T_2, T_3 \rightarrow T_4, T_5 \rightarrow T_1, \dots\}$
 - Example: must put on socks before shoes

Topological Ordering and DAGs

- Want
 - Ordering of tasks which is consistent with dependencies
- Problem representation: directed acyclic graph (DAG)
 - Vertices = tasks; directed edges = dependencies
 - Acyclic: if there exists a cycle of dependencies, no solution possible
- General model for causality, dependency

Topological Ordering Algorithm

TopologicalSort (G)

 Run DFS (G) w/ pre/post numbers


 Return the vertices in *reverse*
 postorder

Note: Can add vertices to list as postorder assigned.

Runtime: $O(|V|+|E|)$.

Topological Ordering

Problem: Design an algorithm that given a DAG G computes a topological ordering on G .



This is just DFS ordering!

Final Ordering: A E B C D

Topological Sort of a DAG

Useful algorithm.

- Linear ordering of vertices s.t. $v \rightarrow w$, implies that v appears before w
- In a DAG, every edge leads to a vertex with lower post number. Why?
 - Any edge (u, v) for which $\text{post}(v) > \text{post}(u)$ is a back edge.
 - Does a DAG have back edges?
- Many graph algorithms are relatively easy to find the answer for v if you've already found the answer for everything downstream of v .
 - Topologically sort G .
 - Solve for v in reverse topological order.

Lecture 5 summary

- Graphs
- Depth First Search
- DAGs and topo sorting